

Module 2.13

Assignment

NTU_CE7

28 Aug 2024

CE7 Group 2 Members

1. Chua Lai Chwang
2. Dijay Kumar
3. Tan Yuan
4. Lovell Tan
5. Khai Razali
6. Shashipal Singh
7. Wong Teck Choy

What is Cloud Architecture
Design – Performance
Efficiency?

Chat generated reference

Cloud architecture design focused on Performance Efficiency, as outlined in the AWS Well-Architected Framework, emphasizes the optimal use of computing resources to meet workload requirements while maintaining efficiency as demands change.

Key Design Principles:

- Select appropriate resource types based on workload requirements
- Leverage managed services to reduce operational overhead
- Embrace serverless architectures for automatic scaling
- Experiment with different configurations to identify optimal performance
- Utilize a data-driven approach using performance metrics

Best Practices:

- Monitor performance continuously using tools like AWS CloudWatch
- Implement auto scaling to dynamically adjust resources based on demand
- Optimize resource allocation by regularly reviewing and adjusting configurations
- Conduct load testing to understand application behavior under various conditions
- Continuously assess and refine architectures based on evolving business needs

Cloud Architecture Design – Performance Efficiency Implementation

How do you select the best performing architecture?

-Tan Yuan

Selection

- **Data-Driven Decisions:** Use a data-driven approach to evaluate and choose different architectural options. The chosen architecture will dictate the AWS services that best optimize performance.
- **Scalability Considerations:** Ensure the architecture can scale efficiently to meet future demand. This may involve selecting services that support auto-scaling or elastic load balancing.

Review

- **Continuous Evaluation:** Regularly review your architecture against emerging technologies like AI and machine learning, which could further enhance performance.
- **Benchmarking:** Conduct benchmarking tests against different architectural options to determine the most effective configuration for your workload.

Monitoring

- **Performance Insights:** Implement continuous monitoring to gather data on system performance and resource utilization. Use these insights to understand and adapt to system-wide changes.
- **Real-Time Alerts:** Set up real-time alerts to quickly identify and address performance bottlenecks or anomalies.

Trade Offs

- **Evaluate Tradeoffs:** Consider the tradeoffs in performance based on your architecture. For example, EBS vs. Instance Store:
- EBS offers persistent storage with additional cost and latency.
- Instance Store provides high-speed, non-persistent storage, suitable for temporary data.
- **Cost vs. Performance:** Balance cost and performance by selecting the most efficient resources that meet your performance needs without overspending.

User Experience

- **Optimize for UX:** Consider the impact of architectural choices on user experience, such as latency, load times, and overall system responsiveness.

How do you monitor your resources to ensure they are performing?

To keep your cloud resources running smoothly and efficiently, consider these best practices:

Utilize Monitoring and Logging Tools

- Native Services: Use integrated tools like AWS CloudWatch for real-time monitoring and logging.
- Enhanced Visibility: Integrate third-party solutions such as Datadog, New Relic, or Prometheus for deeper insights and advanced analytics.

Establish Key Performance Indicators (KPIs)

- Critical Metrics: Focus on essential KPIs like Latency, Throughput, Error Rates, and Resource Utilization to assess performance and health.

Set Up Alerts and Notifications

- Threshold Alerts: Configure alerts to notify you when KPIs exceed or fall below set thresholds.
- Automated Actions: Use notifications to automatically scale resources or prompt manual intervention when necessary.

Analyze Performance Data

- Trend Analysis: Review historical data to uncover performance trends and potential issues.
- Real-Time Monitoring: Keep an eye on live data to quickly identify and resolve performance problems as they arise.

Leverage Automation

- Auto-Scaling: Implement auto-scaling policies to adjust resources based on real-time demand, ensuring efficient utilization.
- Health Checks: Set up automated health checks to regularly verify system stability and performance.

How do you evolve your workload to take advantage of new releases?

1. Stay up-to-date on new resources and services

- Regularly evaluate new AWS services, design patterns, and offerings.
- Engage in ad-hoc evaluations, internal discussions, or external analysis to assess improvements.

1. Define a process to improve workload performance

- Create a process for evaluating new services, designs, and configurations.
- Run performance tests on new AWS offerings to measure potential improvements.

1. Evolve workload performance over time

- Use evaluation data to drive the adoption of new services or resources.
- Adjust to changing performance needs as business or workload demands evolve.

How do you use tradeoffs to improve performance?

- **Understand Workload Requirements:**
 - Analyze the specific needs of your workload to make informed decisions about resource allocation. Doing this can allow for the workload design to cater for the optimum performance as required by business.
- **Select Appropriate Resource Types:**
 - Choose the right instance types and storage solutions that align with performance needs, balancing cost and efficiency. A good practice here will be to “start small and grow bigger” as more performance data-points are gathered and analysis are done.
- **Use Data-Driven Approaches:**
 - Gather and analyze performance metrics to identify bottlenecks and optimize configurations accordingly.
- **Experiment with Different Architectures:**
 - Regularly test various architectures and configurations to find the most efficient setup for your workloads. As much as possible the testing should as automated as possible in order that the risk of changes can be mitigated in production.
- **Implement Caching Strategies:**
 - Utilize caching to reduce latency and improve response times, balancing the trade-off between data freshness and performance.
- **Leverage Auto Scaling:**
 - Use AWS Auto Scaling to dynamically adjust resources based on demand, optimizing performance while controlling costs.
- **Consider Serverless Architectures:**
 - Adopt serverless solutions to eliminate the overhead of managing servers, allowing for automatic scaling and cost efficiency. And to avoid any old-start concern, a pool of pre-allocated resources can be instantiated and ready for use.
- **Optimize Data:**
 - Transfer: Minimize data transfer times by strategically placing resources in geographic locations closer to users. For example, in the case of Web content, AWS Cloudfront can be used to ensure content is as close to the end-user as possible.